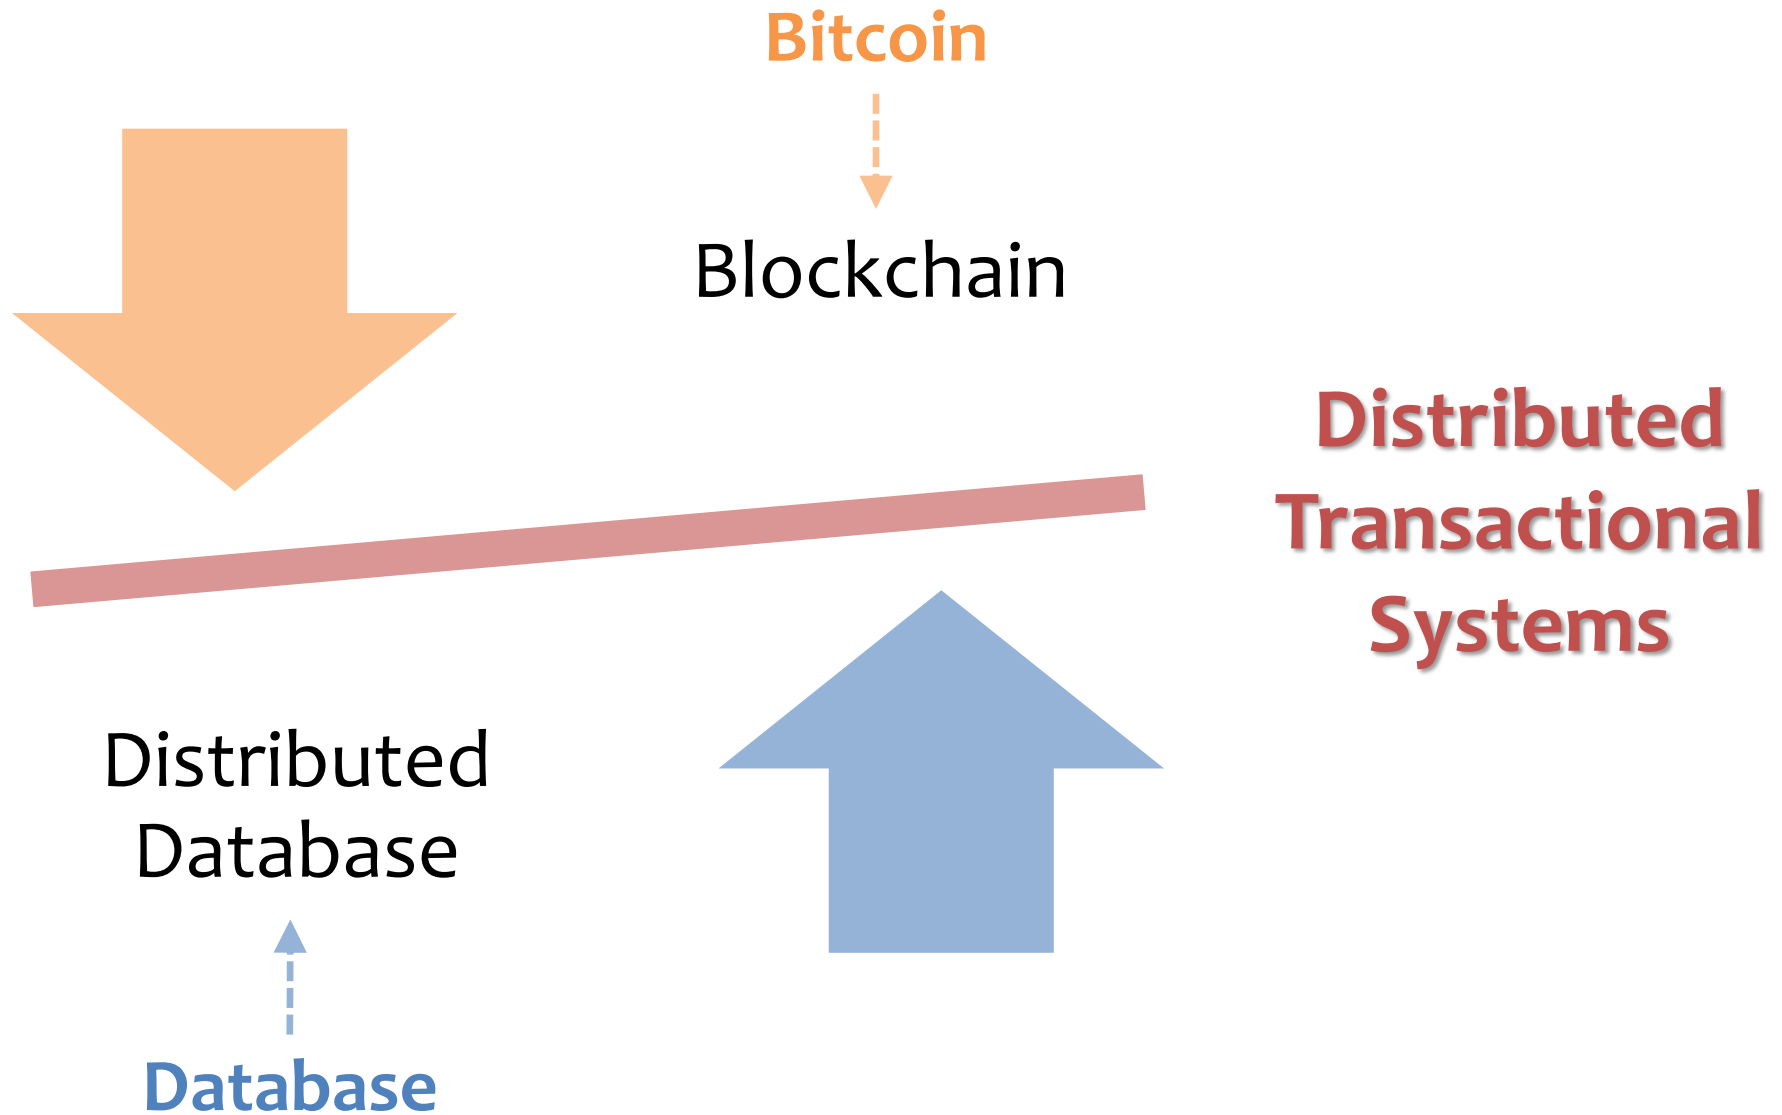


Fine-Grained, Secure and Efficient Data Provenance on Blockchain Systems

Pingcheng RUAN, Gang CHEN, Tien Tuan Anh DINH,
Qian LIN, Beng Chin OOI, Meihui ZHANG




Blockchain Is a Class of Database



Blockchain Basics

- **P2P network**
 - Asynchronous transaction
- **Byzantine environment**
 - Mutual distrusting setup
- **Distributed ledger**
 - Smart contract
- **Inherent provenance-preserving**
 - ONLY for offline analytical query

```
contract Token {  
    method Transfer(sender, recipient, amount) {  
        bal1 = gState[sender];  
        bal2 = gState[recipient];  
        if (amount < bal1) {  
            gState[sender] = bal1 - amount;  
            gState[recipient] = bal2 + amount;  
        }  
    }  
}
```




Motivation

Expose provenance information
to smart contracts both

Efficiently

Securely



**Enabler for provenance-
dependent smart contracts**

Enrich the transaction semantics

Provenance-dependent Contracts

- **Previous transfer precondition:**
 - Enough balance from the sender
 - CURRENT STATE ONLY
- **New transfer precondition:**
 - Historical balance > threshold
 - Recipient not transacted with certain blacklisted addresses recently
 - HISTORICAL STATE & PROVENANCE INFO

```
contract Token {  
    method Transfer(sender, recipient, amount) {  
        bal1 = gState[sender];  
        bal2 = gState[recipient];  
        if (amount < bal1) {  
            gState[sender] = bal1 - amount;  
            gState[recipient] = bal2 + amount;  
        }  
    }  
}
```

Workarounds

Workaround 1:

- Dump every thing into current state
- **Effort-needed, expensive, error-prune**

Workaround 2:

- Offline analytics + Online transactions
- **Break of serializability**
- **Transaction-ordering attacks**

Workaround 3:

- Minimum system instrumentation
- NOT protocol level (e.g., Hyperledger Fabric v1.0+)
- **Data tampering**

Account1_v1: 10
Account1_v2: 20
Account1_v3: 15
Account2_v2: 12

Holistic Approach:

- **Protocol-level enhancement**
→ **Secure**
- **Performance-aware**
→ **Efficient**

Challenges

NO standardized operations

- With clearly-defined transformation semantics
- E.g.
 - Map and reduce in Hadoop
 - Select, join and aggregation in SQL

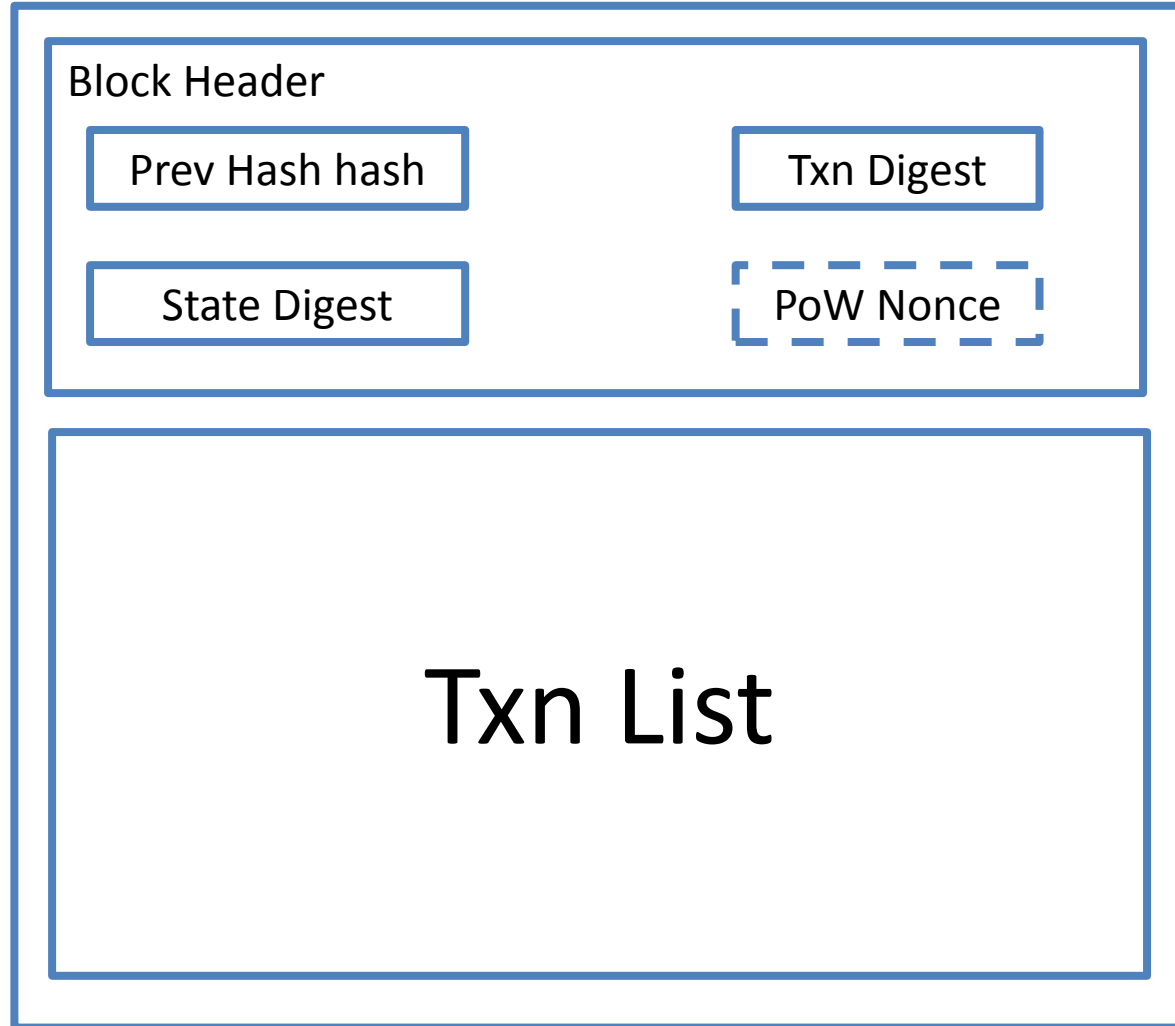
Byzantine environment

- Tamper evidence
- Integrity proof

Ever-growing ledger

- Gas mechanism
- Verifier's dilemma

Block Structure

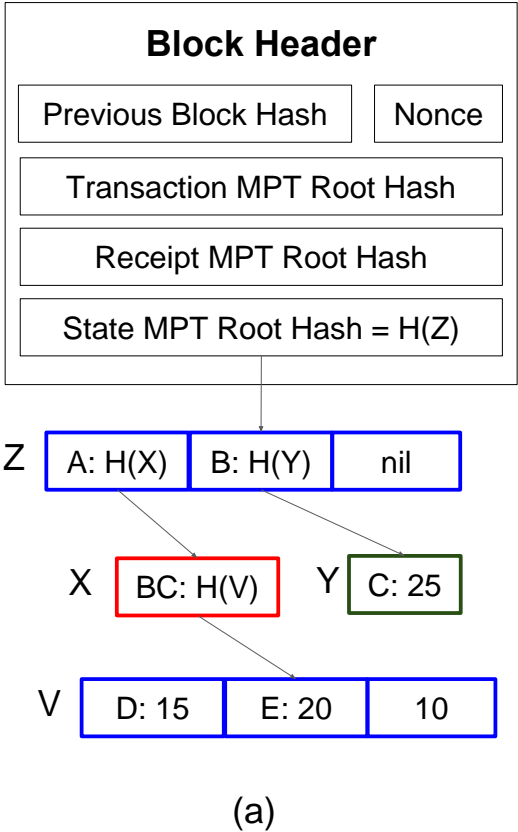


Enhancement Basis (Merkle Tree Variants)

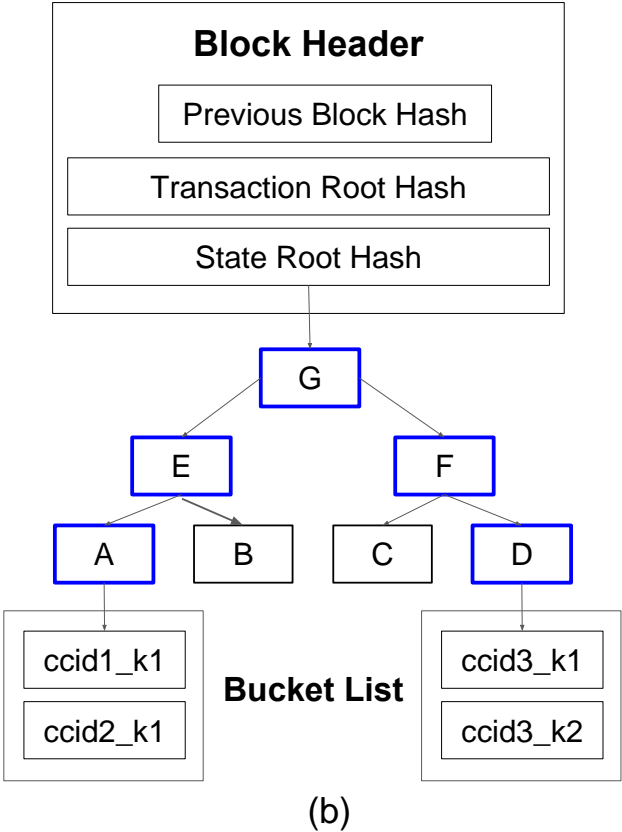
- **Limitation**
 - Latest State only
- **Tamper evidence**
 - Succinct digest (root hash)
 - Integrity proof (access path)

Account Address and Associated Balance in Global State:
0xABC: 10 0xABCD: 15
0xABCE: 20 0xBC: 25

<Updated Chaincode ID>_<Key>:
ccid1_k1 ccid2_k1
ccid3_k1 ccid3_k2

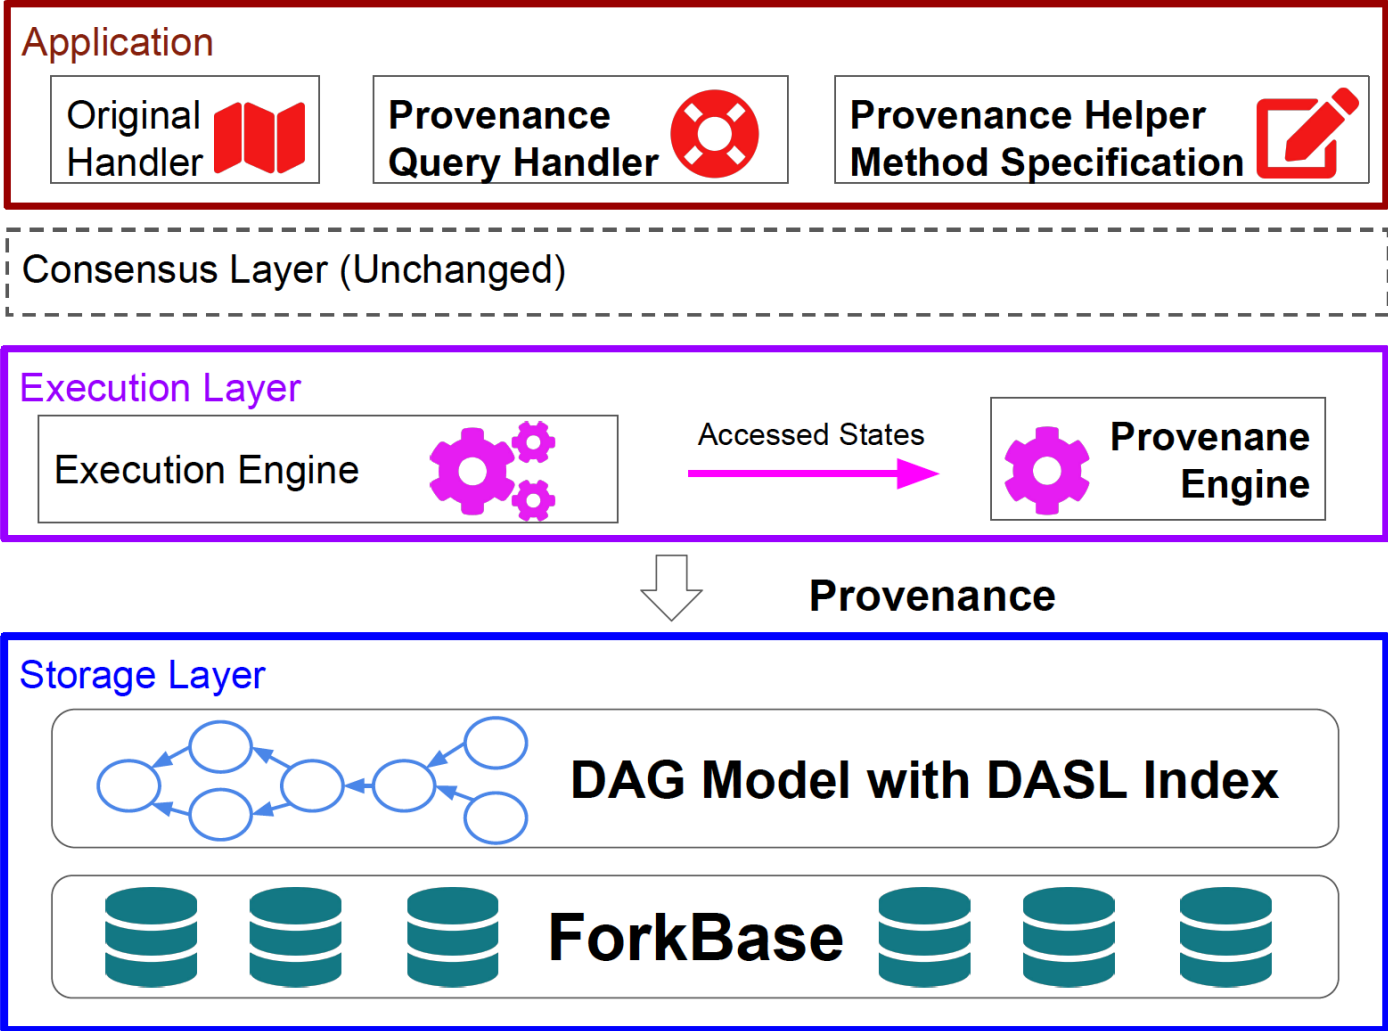


Merkle Patricia Trie



Merkle Bucket Tree

LineageChain Overview






- **Provenance specification**

- User-defined input-output dependency

```
method prov_helper(name: string,  
                   reads: map(string, byte[]),  
                   writes: map(string, byte[]))  
  returns map(string, string[]);
```

- **Provenance query handler**

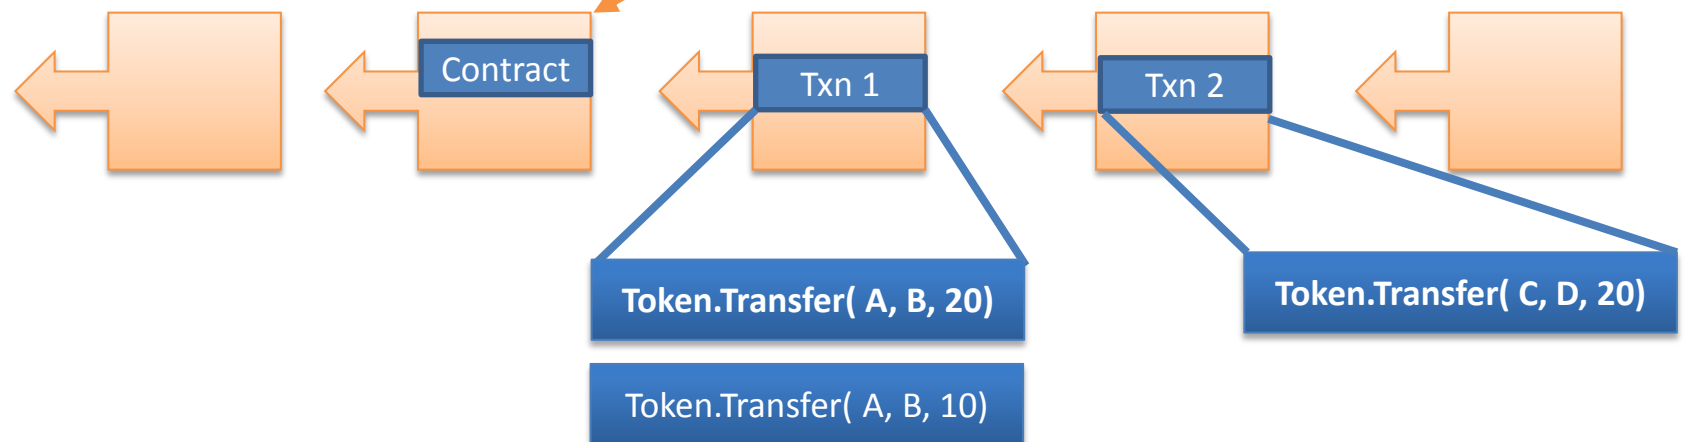
- Hist(stateID, [blockNum])
→ (val, blkStart, txnID)
- Backward(stateID, blkNum)
→ List<(depStateID, depBlkNum)>
- Forward(stateID, blkNum)
→ List<(depStateID, depBlkNum)>



Blockchain Basics

- **P2P network**
 - Asynchronous transaction
- **Byzantine environment**
 - Mutual distrusting setup
- **Distributed ledger**
 - Smart contract
- **Inherent provenance-preserving**
 - ONLY for offline analytical query

```
contract Token {  
    method Transfer(sender, recipient, amount) {  
        bal1 = gState[sender];  
        bal2 = gState[recipient];  
        if (amount < bal1) {  
            gState[sender] = bal1 - amount;  
            gState[recipient] = bal2 + amount;  
        }  
    }  
}
```






- **Provenance specification**

- User-defined input-output dependency

```
method prov_helper(name: string,  
                   reads: map(string, byte[]),  
                   writes: map(string, byte[]))  
  returns map(string, string[]);
```

- **Provenance query handler**

- Hist(stateID, [blockNum])
→ (val, blkStart, txnID)
- Backward(stateID, blkNum)
→ List<(depStateID, depBlkNum)>
- Forward(stateID, blkNum)
→ List<(depStateID, depBlkNum)>



Application Layer

Application

Original
Handler



Provenance
Query Handler



Provenance Helper
Method Specification



```
contract Token {
  method Transfer(...){...} // as above
  method prov_helper(name, reads, writes) {
    if name == "Transfer" {
      for (id,value) in writes {
        if (reads[id] < value) {
          recipient = id;
        } else {sender = id; }
      }
      // dependency list with a
      // single element.
      dep = [sender];
      return {recipient:dep};
    }
    ...
  }
}

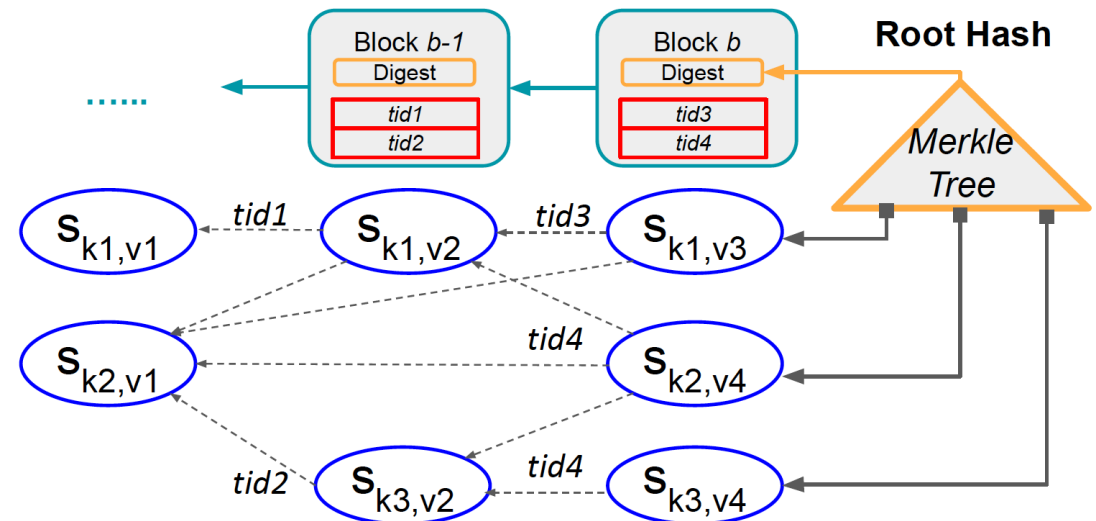


method Refund(addr) {
  blk := last block in the ledger
  first_blk := first block in this month
  sum = count = 0;
  while (first_blk < blk) {
    val, startBlk, txnID = Hist(addr, blk);
    blk = startBlk - 1;
    sum += val;
    count += 1;
  }
  avg = sum / count;
  refund_amount := refund amount based on avg
  gState[addr] += refund_amount;
}
```

Recipient -> Sender

```
method Blacklist(addr) {
  blk := last block in the ledger
  blacklisted = false;
  iterate 5 times {
    val, startBlk, txnID = Hist(addr, blk);
    for (depAddr, depBlk)
      in (Backward(addr, startBlk)
         or Forward(addr, startBlk)) {
      if depAddr in gState["blacklist"] {
        gState["blacklist"].append(addr);
        return;
      }
    }
    blk = startBlk - 1;
  }
}
```

Execution Layer

- **Receive**
 - Contract invocation context
 - Provenance specification
- **Compute**
 - Transaction results
 - Concrete dependency
- **Prepare Merkle DAG**
 - Introduce one layer of direction
 - Hash reference to encode provenance backward dependency



Execution Layer

Execution Layer

Execution Engine




Accessed States





Provenance Engine

- **Forward tracking**
 - Problem: Undecided forward dependency during state update
- **Solution**
 - Lazily store forward dependency on the successor state entry

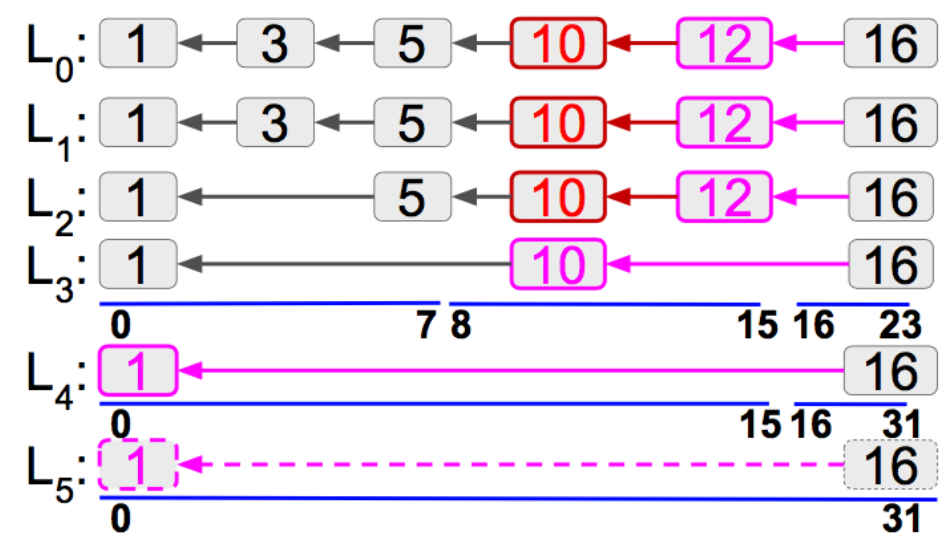


Storage Layer

- **Problem**
 - Efficient version-based (historical) query for a state ID
- **Solution:**
 - Deterministic Append-only Skip List
 - Hash-based reference



After appending
versions 12 and 16



DEFINITION 6. Let $V_k = \langle v_0, v_1, \dots \rangle$ be the sequence of version numbers of states with identifier k , in which $v_i < v_j$ for all $i < j$. A DASL index for k consists of N linked lists L_0, L_1, \dots, L_{N-1} . Let v_j^{i-1} and v_j^i be the versions in the $(j-1)^{th}$ and j^{th} node of list L_i . Let b be the base number, a system-wide parameter. The content of L_i is constructed as follows:

- 1) $v_0 \in L_i$
- 2) Given v_{j-1}^i , v_j^i is the smallest version in V_k such that:

$$\left\lfloor \frac{v_{j-1}^i}{b^i} \right\rfloor < \left\lfloor \frac{v_j^i}{b^i} \right\rfloor \quad (5)$$


Evaluation

- **MICRO benchmarking** (vs. flat storage)
 - Preference to recent version query (with DASL)
 - More efficient BFS enabled by backtrack (with ForkBase)
- **MACRO benchmarking** (applied to Hyperledger Fabric v0.6 and v1.3)
 - Negligible runtime overhead
 - Tiny proportion of latency
 - Negligible storage overhead
 - >70% of space for blocks
 - 25% for historical states
 - 2~4% for DASL indexes and hash pointers

Performance of Provenance Query




- **vs. Workaround 2**

- Compute data provenance offline and conditionally trigger online transaction




Micro Performance of Provenance Query

- **vs. Workaround 1**
 - Dump everything into the current state
- **vs. Workaround 3**
 - Use Hyperledger Fabric's built-in HistoryDB




Runtime Overhead

- Transaction processing




Hyperledger Fabric v0.6



Hyperledger Fabric v1.3

Storage Overhead



Conclusion

- **LineageChain**
 - Enabler for provenance-dependent blockchain applications
 - Protocol-level enhancement w.r.t. efficiency and security
 - Negligible performance and storage overhead
- **Key designs**
 - User-defined dependency specification
 - Merkle DAG with dependency tracking
 - DASL index to accelerate data provenance query
 - Adoption in Hyperledger Fabric (v0.6 & v1.3)

Thank You!